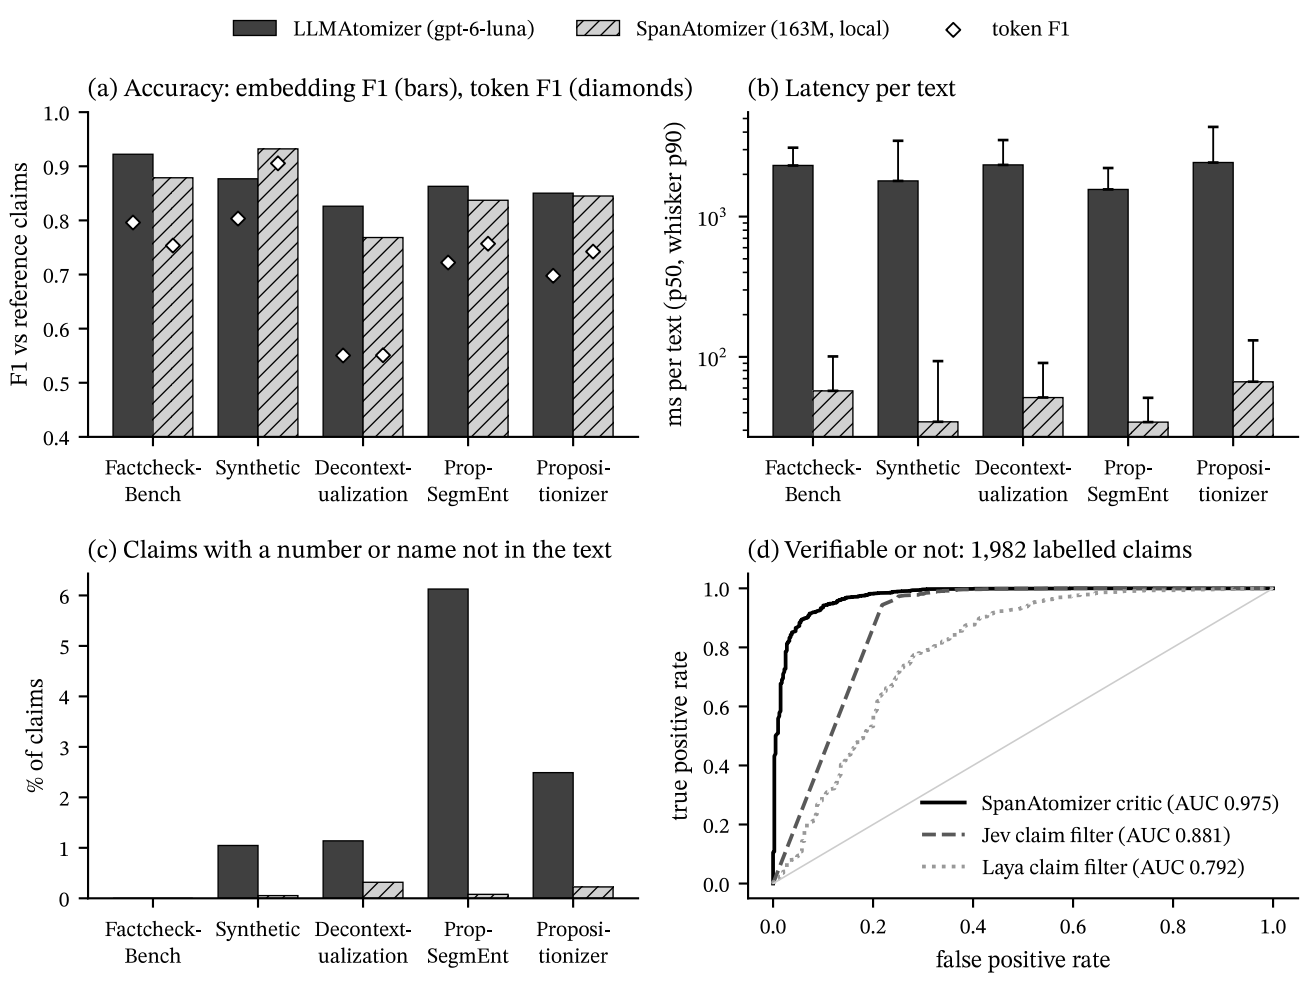

LLMAtomizer (gpt-6-luna)
SpanAtomizer (163M, local)
token F1
(a) Accuracy: embedding F1 (bars), token F1 (diamonds)
F1 vs reference claims
1.0
0.9
0.8
0.7
0.6
0.5
0.4
Factcheck-Bench
Synthetic
Decontext-ualization
Prop-SegmEnt
Proposi-tionizer
(b) Latency per text
ms per text (p50, whisker p90)
10^3
10^2
Factcheck-Bench
Synthetic
Decontext-ualization
Prop-SegmEnt
Proposi-tionizer
(c) Claims with a number or name not in the text
% of claims
6
5
4
3
2
1
0
Factcheck-Bench
Synthetic
Decontext-ualization
Prop-SegmEnt
Proposi-tionizer
(d) Verifiable or not: 1,982 labelled claims
true positive rate
1.0
0.8
0.6
0.4
0.2
0.0
SpanAtomizer critic (AUC 0.975)
Jev claim filter (AUC 0.881)
Laya claim filter (AUC 0.792)
0.0
0.2
0.4
0.6
0.8
1.0
false positive rate